

### **3.1**

How long (in days and hours) clktimeilli can keep time after bootloading until the counter overflows. Truncate fractional hours. For how long can a 64-bit millisecond counter keep time?

**As a uint32 variable can have a value of  $2^{32}-1$ . So clktimeilli has a maximum value of  $2^{32}-1$ . Converting this to be representative of days and hours is as follows:**

**$2^{32}-1 / 1000 / 60 / 60 / 24 = 49$  Days**

**$(2^{32}-1 \% (1000 * 60 * 60 * 24)) / (1000 * 60 * 60) = 17$  hours**

**So, clktimeilli can record up to 49 days 17 hours.**

**A 64-bit millisecond counter would be able to keep time for 213503982334 days 14 hours.**

**OR 584942417 years 14 hours.**

### **3.4**

Explain in Lab3Answers.pdf why adding instrumentation code before and after calling `ctxsw()` in `resched()` is a viable method for estimating a process's gross CPU usage. Describe the sequence of XINU kernel function calls that are evoked in the three scenarios. Explain why the 3.3 method will perform correct gross CPU usage tracking in the three scenarios.

**This is a viable way of estimating a process's gross CPU usage because the amount of time that the new process is switched in is calculated by the difference between `clktimemilli` before and after `ctxsw()` is called. This essentially counts the time in milliseconds that the new process was being run.**

Scenario 1: the current process makes a `sleepms()` system call which context-switches out the current process and context-switches in a new (i.e., different) process.

**`sleepms()` disables interrupts, changes process state to sleep, calls `insertd()`**

**`insertd()` inserts the process into the delta list, returns back to `sleepms()`**

**`sleepms()` calls `resched()`**

**`resched()` selects the next process in ready queue**

**The clock timers start.**

**`resched()` calls `ctxsw()`**

**`ctxsw()` switches the old processes stack with the new process stack, restores interrupts, returns to new process**

**Eventually, `resched()` is called in the new process, the old process is resumed in the call to `ctxsw()` which then returns back to `resched()`.**

**Then the gross times are updated.**

Scenario 2: while the current process is executing, a clock interrupt is raised which causes the current process's time slice to be decremented by `clkhandler()`. If the remaining time slice becomes 0, `resched()` is called which may, or may not, trigger a context-switch depending on the priority of the process at the front of XINU's ready list.

**Old process is interrupted by a clock interrupt, `clkhandler()` is called**

**`clkhandler()` decrements preempt, so `resched()` is called.**

**`resched()` selects a new process to run from the ready list and calls `ctxsw()`; however, if the process selected has a lower priority than the old process, then the old process will resume with a fresh preempt.**

**The clock timers start.**

**If `ctxsw()` is called, then it switches the old processes stack with the new process stack and resumes the new process**

**Eventually, `resched()` is called in the new process, the old process is resumed in the call to `ctxsw()` which then returns back to `resched()`.**

**Then the gross times are updated.**

Scenario 3: while the current process is executing, a clock interrupt is raised which causes a previously sleeping process to be woken up and placed into the ready list. If the newly awoken process has priority greater or equal than the current process, a context-switch ensues.

**Old process is interrupted by a clock interrupt, clkhandler() is called**

**clkhandler() checks to see if the delta list (sleep list) is empty. If it isn't and the process is done waiting, then call wakeup()**

**wakeup() calls resched\_cntl(DEFER\_START)**

**resched\_cntl() increments ndefers and returns back to wakeup()**

**wakeup() then turns all of the sleeping processes with no more time waiting to ready. Then it calls resched\_cntl(DEFER\_STOP)**

**resched\_cntl() decrements ndefers then calls resched()**

**resched() now selects a new process to run from the ready list. Now assuming that the newly awoken process has the highest priority, then this process is chosen as the new process. ctxsw() is called.**

**The clock timers start.**

**ctxsw() it switches the old processes stack with the new process stack and resumes the new process**

**Eventually, resched() is called in the new process, the old process is resumed in the call to ctxsw() which then returns back to resched().**

**Then the gross times are updated.**

### **3.5**

Perform benchmark runs and compare the values returned by `procgrosscpumicro()` against those returned by `procgrosscpu()`.

**I created a process with an infinite loop in it, so that the gross cpu time will increase over time. The value returned by `procgrosscpu()` \* 1000 will always be higher than the value returned by `procgrosscpumicro()` because `resched()` rounds up the gross cpu time whereas the gross cpu time in microseconds is dependent on the number of cpu ticks, which is never rounded. With this being said, the longer a process is active, the more inaccurate `procgrosscpu()` will become.**

### **4.2**

Describe changes made to support non-decreasing priority system:

**Changes made to support non-decreasing priority system:**

- **Added MAXPRIO to process.h**
- **Changed NULLPROC to have priority MAXPRIO**
- **Changed empty proc table entries to have prio MAXPRIO**
- **Changed the ready list so that the head has MINKEY and the tail has MAXKEY to account for the non-decreasing order.**
- **Changed the checks in `rcreate` to accomodate for non-decreasing**
- **Replaced calls to `insert` with `rinsert`.**
- **Changed comparator in `resched()` to accomodate for non-decreasing**

### **4.3**

The null process must be treated as a special case so that its CPU usage is strictly greater than all other processes. This implies that the null process, when not current, is always at the end of the ready list. Explain in Lab3Answers.pdf how you ensure that this is the case.

**I am ensuring that the NULLPROC always has the highest value of prvgrosscpu so that it is guaranteed to be the last item in the readylist. Wherever the values of prvgrosscpu are updated, a check is made to ensure that the NULLPROC's prvgrosscpu value will always be at least one value higher. Also, I added checks to physically check to see if NULLPROC is the last item in ready list, if it is not, then rearrange the ready list to make it happen.**

### **4.4**

XINU's null process (PID 0) must be treated separately so that its priority is always strictly less than the priority of all other processes in the system. Explain in Lab3Answers.pdf how you go about assuring that.

**Since we are using the prvgrosscpu value as the priority, a higher priority means a lower prvgrosscpu value and vice versa. Therefore, in order to ensure that the NULL process has the lowest priority, we have to ensure that the value of prvgrosscpu time is the highest at all times. This is taken care of by 4.3.**

### **5.2**

Create 4 CPU-bound processes from main() back-to-back. If your scheduler is implemented correctly, we would expect to see the 4 processes printing similar CPU usage and x values. Repeat the benchmark test one more time and inspect your results.

**With my implementation, this is the output for trials 1 and 2:**

cpu: 3	124143507	4051	4039424	cpu: 3	124145019	4051	4039429
cpu: 4	123779488	3961	3950290	cpu: 4	123779171	3961	3950290
cpu: 5	123219319	3961	3950289	cpu: 5	123219470	3961	3950288
cpu: 6	124689847	3963	3950284	cpu: 6	124690483	3963	3950284

**With my implementation, I can see that the cpu is shared between the processes in a pretty equal way.**

### 5.3

Create 4 I/O-bound processes from main() and perform the same benchmark tests as 5.2.

With my implementation, this is the output for trials 1 and 2:

io: 7	160	160	132507	io: 7	160	160	132505
io: 8	160	7951	7886705	io: 8	160	7951	7886706
io: 10	160	209	93706	io: 10	160	209	93706
io: 9	160	7951	7887891	io: 9	160	7951	7887892

I see that the value of x is the exact same throughout the four tests, and two processes have a very similar cputime and the other two processes have a very similar cputime. I was not able to figure out why the io devices did not have better sharing.

### 5.4

Create 4 CPU-bound processes and 4 I/O-bound processes. We would expect the 4 CPU-bound processes to output similar x values and CPU usage with respect to each other, and the same goes for the 4 I/O-bound processes within the group of I/O-bound processes. Across the two groups, we would expect CPU-bound processes to receive significantly more CPU time than I/O-bound processes

**Trials 1 and 2:**

cpu: 11	124163986	4051	4040034	cpu: 11	124163960	4051	4040033
cpu: 12	123776656	3961	3950288	cpu: 12	123776689	3961	3950289
cpu: 13	124135331	4021	4010132	cpu: 13	124135239	4021	4010133
cpu: 14	125556296	4033	4020115	cpu: 14	125556303	4033	4020115
io: 15	120	7979	7925947	io: 15	120	7979	7925943
io: 17	119	2112	2104680	io: 17	119	2112	2104680
io: 16	120	8008	7954864	io: 16	120	8008	7954864
io: 18	119	3461	3466160	io: 18	119	3461	3466170

In this test, I see that the IO test produce about same count of x, and have relatively similar grosscputime throughout the tests. The cpu tests are very similar and have good sharing of the CPU.

## 5.5

A variant of test scenario A, create 4 CPU-bound processes in sequence with 500 msec delays (by calling `sleepms()`) added between successive `create()` (nested with `resume()`) system calls. Estimate how much CPU time the first process should receive. The same goes for the other 3 processes. Compare your calculations with the actual performance results from testing.

### **Trials 1 and 2:**

cpu: 19	129987376	8556	8512063	cpu: 19	129986957	8556	8512069
cpu: 20	122306299	7521	7495588	cpu: 20	122305593	7521	7495576
cpu: 21	123226062	3961	3950283	cpu: 21	123226060	3961	3950283
cpu: 22	131135776	3767	3740803	cpu: 22	131135781	3767	3740803

**Adding sleep calls in between creation of processes yields interesting results. The first two and the last two processes have similar cpu times, but the value of x is the same throughout.**